

TELECOM CUSTOMER CHURN PREDICTION

Yipeng CHEN



Figure 1: Customer Churn

Did you know that attracting a new customer costs five times as much as keeping an existing one?

Initially,

I would like to express my gratitude to the following Kaggle notebooks that have served as a source of inspiration for the creation of this report:

[CUSTOMER CHURN PREDICTION - bhartiprasad17](#)

[Telecom Churn Prediction - bandiatindra](#)

Contents

1. Introduction	3
1.1 About Dataset	4
2. Loading libraries and data	5
3. Understanding the data	6
4. Visualize missing values	9
5. Data Manipulation	11
6. Data Visualization	14
7. Data Preprocessing	35
8. Machine Learning Model Evaluations and Predictions	41
8.1 Logistic Regression	41
8.2 AdaBoost	41
8.3 Gradient Boosting	42
8.4 Voting Classifier	42
8.5 Feature importance based on Voting Classifier Model	43
9. Advice	45

1. Introduction

What is Customer Churn?

Customer churn or customer attrition is defined as when customers discontinue using a company's product or service.

Individualized customer retention is tough because most firms have a large number of customers and can't afford to devote much time to each of them. The costs would be too great, outweighing the additional revenue. However, if a corporation could forecast which customers are likely to leave ahead of time, it could focus customer retention efforts only on these "high risk" clients. The ultimate goal is to expand its coverage area and retrieve more customers loyalty.

Customer churn is a critical metric because it is much less expensive to retain existing customers than it is to acquire new customers. Moreover, customer churn is a giant business killer. Even small increases in churn can cause a significant cut in revenues.

To reduce customer churn, companies need to predict which customers are at high risk of churn.

To detect early signs of potential churn, one must first develop a holistic view of the customers and their interactions across numerous channels, including store/branch visits, product purchase histories, customer service calls, web-based transactions, and social media interactions, to mention a few.

As a result, by addressing churn, these businesses may not only preserve their market position, but also grow and thrive. More customers they have in their network, the lower the cost of initiation and the larger the profit. As a result, the company's key focus for success is reducing client attrition and implementing effective retention strategy.

1.1 About Dataset

Context

IBM Sample Data Sets

“Predict behavior to retain customers. You can analyze all relevant customer data and develop focused customer retention programs.”

Content

Each row represents a customer, each column contains customer’s attributes.

The data set includes information about:

- Customers who left within the last month – the column is called Churn
- Services that each customer has signed up for – phone, multiple lines, internet, online security, online backup, device protection, tech support, and streaming TV and movies
- Customer account information – how long they’ve been a customer, contract, payment method, paperless billing, monthly charges, and total charges
- Demographic info about customers – gender, age range, and if they have partners and dependents

The raw data contains 7043 rows (customers) and 21 columns (features).

The “Churn” column is our target.

Columns description

Column_Name	Column_Description
customerID	Customer ID
gender	Whether the customer is a male or a female
SeniorCitizen	Whether the customer is a senior citizen or not (1, 0)
Partner	Whether the customer has a partner or not (Yes, No)
Dependents	Whether the customer has dependents or not (Yes, No)
tenure	Number of months the customer has stayed with the company
PhoneService	Whether the customer has a phone service or not (Yes, No)
MultipleLines	Whether the customer has multiple lines or not (Yes, No, No phone service)
InternetService	Customer’s internet service provider (DSL, Fiber optic, No)
OnlineSecurity	Whether the customer has online security or not (Yes, No, No internet service)
OnlineBackup	Whether the customer has online backup or not (Yes, No, No internet service)
DeviceProtection	Whether the customer has device protection or not (Yes, No, No internet service)
TechSupport	Whether the customer has tech support or not (Yes, No, No internet service)
StreamingTV	Whether the customer has streaming TV or not (Yes, No, No internet service)
StreamingMovies	Whether the customer has streaming movies or not (Yes, No, No internet service)
Contract	The contract term of the customer (Month-to-month, One year, Two year)
PaperlessBilling	Whether the customer has paperless billing or not (Yes, No)
PaymentMethod	The customer’s payment method (Electronic check, Mailed check, Bank transfer (automatic), Credit card (automatic))
MonthlyCharges	The amount charged to the customer monthly
TotalCharges	The total amount charged to the customer
Churn	Whether the customer churned or not (Yes or No)

2. Loading libraries and data

```
import pandas as pd
import numpy as np
import missingno as msno
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import plotly.graph_objects as go
from plotly.subplots import make_subplots
import warnings
warnings.filterwarnings('ignore')

from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import GradientBoostingClassifier

from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

# Loading data
df = pd.read_csv('Telco-Customer-Churn.csv')
```

3. Understanding the data

```
df
```

```
##      customerID  gender  SeniorCitizen  ...  MonthlyCharges  TotalCharges  Churn
## 0      7590-VHVEG  Female                0  ...           29.85           29.85    No
## 1      5575-GNVDE   Male                0  ...           56.95          1889.5    No
## 2      3668-QPYBK   Male                0  ...           53.85           108.15   Yes
## 3      7795-CFOCW   Male                0  ...           42.30          1840.75    No
## 4      9237-HQITU  Female                0  ...           70.70           151.65   Yes
## ...      ...      ...
## 7038  6840-RESVB   Male                0  ...           84.80          1990.5    No
## 7039  2234-XADUH  Female                0  ...          103.20          7362.9    No
## 7040  4801-JZAZL  Female                0  ...           29.60           346.45    No
## 7041  8361-LTMKD   Male                1  ...           74.40            306.6   Yes
## 7042  3186-AJIEK   Male                0  ...          105.65          6844.5    No
##
## [7043 rows x 21 columns]
```

```
df.shape
```

```
## (7043, 21)
```

```
df.columns.values
```

```
## array(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
##       'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
##       'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
##       'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract',
##       'PaperlessBilling', 'PaymentMethod', 'MonthlyCharges',
##       'TotalCharges', 'Churn'], dtype=object)
```

```
df.info()
```

```
## <class 'pandas.core.frame.DataFrame'>
## RangeIndex: 7043 entries, 0 to 7042
## Data columns (total 21 columns):
##  #   Column                Non-Null Count  Dtype
## ---  ---
##  0   customerID            7043 non-null   object
##  1   gender                7043 non-null   object
##  2   SeniorCitizen         7043 non-null   int64
##  3   Partner               7043 non-null   object
##  4   Dependents            7043 non-null   object
##  5   tenure                7043 non-null   int64
##  6   PhoneService          7043 non-null   object
##  7   MultipleLines         7043 non-null   object
##  8   InternetService       7043 non-null   object
##  9   OnlineSecurity        7043 non-null   object
## 10   OnlineBackup          7043 non-null   object
## 11   DeviceProtection      7043 non-null   object
## 12   TechSupport           7043 non-null   object
## 13   StreamingTV           7043 non-null   object
## 14   StreamingMovies       7043 non-null   object
## 15   Contract              7043 non-null   object
## 16   PaperlessBilling      7043 non-null   object
## 17   PaymentMethod         7043 non-null   object
## 18   MonthlyCharges        7043 non-null   float64
## 19   TotalCharges          7043 non-null   object
## 20   Churn                 7043 non-null   object
## dtypes: float64(1), int64(2), object(18)
## memory usage: 1.1+ MB
```

```
df.dtypes
```


```
## customerID      object
## gender           object
## SeniorCitizen    int64
## Partner          object
## Dependents       object
## tenure           int64
## PhoneService     object
## MultipleLines    object
## InternetService  object
## OnlineSecurity   object
## OnlineBackup     object
## DeviceProtection object
## TechSupport      object
## StreamingTV      object
## StreamingMovies  object
## Contract         object
## PaperlessBilling object
## PaymentMethod    object
## MonthlyCharges   float64
## TotalCharges     object
## Churn            object
## dtype: object
```

```
# Converting "TotalCharges" to a numerical data type
```

```
df['TotalCharges'] = pd.to_numeric(df.TotalCharges, errors='coerce')
```


4. Visualize missing values

```
# Visualize missing values as a bar plot  
msno.bar(df,fontsize=11,figsize=(20,10))
```



From the above visualization,
we can observe that there are 11 missing values for “TotalCharges”.

```
# Calculate the number of missing values in each column of the DataFrame  
df.isnull().sum()
```

```
## customerID          0  
## gender              0  
## SeniorCitizen      0  
## Partner             0  
## Dependents         0  
## tenure             0  
## PhoneService       0  
## MultipleLines      0  
## InternetService    0  
## OnlineSecurity     0  
## OnlineBackup       0  
## DeviceProtection   0  
## TechSupport        0  
## StreamingTV        0  
## StreamingMovies    0  
## Contract           0  
## PaperlessBilling   0  
## PaymentMethod      0  
## MonthlyCharges     0  
## TotalCharges       11  
## Churn              0  
## dtype: int64
```

5. Data Manipulation

```
# Remove customer IDs from the data set
df = df.drop(['customerID'], axis = 1)
df
```

```
##      gender SeniorCitizen Partner ... MonthlyCharges TotalCharges Churn
## 0      Female              0     Yes ...         29.85         29.85    No
## 1       Male              0     No  ...         56.95        1889.50    No
## 2       Male              0     No  ...         53.85         108.15   Yes
## 3       Male              0     No  ...         42.30        1840.75    No
## 4      Female              0     No  ...         70.70         151.65   Yes
## ...      ...              ...     ... ...          ...          ...    ...
## 7038     Male              0     Yes ...         84.80        1990.50    No
## 7039  Female              0     Yes ...        103.20        7362.90    No
## 7040  Female              0     Yes ...         29.60         346.45    No
## 7041     Male              1     Yes ...         74.40         306.60   Yes
## 7042     Male              0     No  ...        105.65        6844.50    No
##
## [7043 rows x 20 columns]
```

```
# We know that the "TotalCharges" has 11 missing values. Let's check this :
df[np.isnan(df['TotalCharges'])]
```

```
##      gender SeniorCitizen Partner ... MonthlyCharges TotalCharges Churn
## 488  Female              0     Yes ...         52.55          NaN    No
## 753   Male              0     No  ...         20.25          NaN    No
## 936  Female              0     Yes ...         80.85          NaN    No
## 1082   Male              0     Yes ...         25.75          NaN    No
## 1340  Female              0     Yes ...         56.05          NaN    No
## 3331   Male              0     Yes ...         19.85          NaN    No
## 3826   Male              0     Yes ...         25.35          NaN    No
## 4380  Female              0     Yes ...         20.00          NaN    No
## 5218   Male              0     Yes ...         19.70          NaN    No
## 6670  Female              0     Yes ...         73.35          NaN    No
## 6754   Male              0     No  ...         61.90          NaN    No
##
## [11 rows x 20 columns]
```

```
# Now, Let us remove these 11 rows from our data set :  
df.dropna(inplace = True)  
df.isnull().sum()
```

```
## gender          0  
## SeniorCitizen  0  
## Partner         0  
## Dependents     0  
## tenure         0  
## PhoneService   0  
## MultipleLines   0  
## InternetService 0  
## OnlineSecurity  0  
## OnlineBackup    0  
## DeviceProtection 0  
## TechSupport     0  
## StreamingTV     0  
## StreamingMovies 0  
## Contract        0  
## PaperlessBilling 0  
## PaymentMethod   0  
## MonthlyCharges  0  
## TotalCharges    0  
## Churn           0  
## dtype: int64
```

```
df.shape
```

```
## (7032, 20)
```

```
df["SeniorCitizen"] = df["SeniorCitizen"].map({0: "No", 1: "Yes"})
df.head()
```

```
##      gender SeniorCitizen Partner  ... MonthlyCharges  TotalCharges  Churn
## 0  Female           No      Yes  ...         29.85         29.85    No
## 1   Male           No      No   ...         56.95        1889.50    No
## 2   Male           No      No   ...         53.85         108.15   Yes
## 3   Male           No      No   ...         42.30        1840.75    No
## 4  Female           No      No   ...         70.70         151.65   Yes
##
## [5 rows x 20 columns]
```

```
df["InternetService"].describe(include=['object', 'bool'])
```

```
## count          7032
## unique           3
## top      Fiber optic
## freq          3096
## Name: InternetService, dtype: object
```

```
numerical_cols = ['tenure', 'MonthlyCharges', 'TotalCharges']
df[numerical_cols].describe()
```

```
##      tenure  MonthlyCharges  TotalCharges
## count  7032.000000    7032.000000    7032.000000
## mean    32.421786     64.798208     2283.300441
## std     24.545260     30.085974     2266.771362
## min      1.000000     18.250000     18.800000
## 25%      9.000000     35.587500     401.450000
## 50%     29.000000     70.350000    1397.475000
## 75%     55.000000     89.862500    3794.737500
## max     72.000000    118.750000    8684.800000
```

6. Data Visualization

```
# g_labels = ['Male', 'Female']
# c_labels = ['No', 'Yes']
#
# # Create subplots: use 'domain' type for Pie subplot
# fig = make_subplots(rows=1, cols=2,
#                     specs=[[{'type': 'domain'}, {'type': 'domain'}]])
# fig.add_trace(go.Pie(labels=g_labels,
#                     values=df['gender'].value_counts(), name="Gender"),
#             1, 1)
# fig.add_trace(go.Pie(labels=c_labels,
#                     values=df['Churn'].value_counts(), name="Churn"),
#             1, 2)
#
# # Use `hole` to create a donut-like pie chart
# fig.update_traces(hole=.4, hoverinfo="label+percent+name", textfont_size=16)
#
# fig.update_layout(
#     title_text="Gender and Churn Distributions",
#     # Add annotations in the center of the donut pies.
#     annotations=[dict(text='Gender', x=0.16, y=0.5,
#                     font_size=20, showarrow=False),
#                 dict(text='Churn', x=0.84, y=0.5,
#                     font_size=20, showarrow=False)])
# fig.show()
```

Gender and Churn Distributions




Figure 2: Gender and Churn Distributions

26.6 % of customers switched to another firm.

Customers are 49.5 % female and 50.5 % male.

```
df["Churn"][df["Churn"]=="No"].groupby(by=df["gender"]).count()
```

```
## gender
## Female    2544
## Male      2619
## Name: Churn, dtype: int64
```

```
df["Churn"][df["Churn"]=="Yes"].groupby(by=df["gender"]).count()
```

```
## gender
## Female     939
## Male       930
## Name: Churn, dtype: int64
```

```

plt.figure(figsize=(6, 6))
labels = ["Churn: Yes", "Churn:No"]
values = [1869, 5163]
labels_gender = ["F", "M", "F", "M"]
sizes_gender = [939, 930, 2544, 2619]
colors = ['#ff6666', '#66b3ff']
colors_gender = ['#c2c2f0', '#ffb3e6', '#c2c2f0', '#ffb3e6']
explode = (0.3, 0.3)
explode_gender = (0.1, 0.1, 0.1, 0.1)
textprops = {"fontsize": 15}
#Plot
plt.pie(values, labels=labels, autopct='%1.1f%%',
pctdistance=1.08, labeldistance=0.8, colors=colors, startangle=90, frame=True,
explode=explode, radius=10, textprops = textprops, counterclock = True, )

```

```

plt.pie(sizes_gender, labels=labels_gender,
colors=colors_gender, startangle=90, explode=explode_gender, radius=7,
textprops = textprops, counterclock = True, )
#Draw circle

```

```

centre_circle = plt.Circle((0,0), 5, color='black', fc='white', linewidth=0)
fig = plt.gcf()
fig.gca().add_artist(centre_circle)

```

```

plt.title('Churn Distribution w.r.t Gender: Male(M), Female(F)',
fontsize=15, y=1.1)

```

```

# show plot
plt.axis('equal')


```

```

plt.tight_layout()
plt.show()

```


Churn Distribution w.r.t Gender: Male(M), Female(F)



There is negligible difference in customer percentage.
Both genders behaved in similar fashion when it comes to migrating to another service provider.

```
# fig = px.histogram(df, x="Churn", color="Contract", barmode="group",
#                    title="Customer contract distribution")
# fig.update_layout(width=700, height=500, bargap=0.1)
# fig.show()
```




Figure 3: Customer contract distribution

About 88% of customers with Month-to-Month Contract opted to move out compared to 9% of customers with One Year Contract and 3% with Two Year Contract.

```
# labels = df['PaymentMethod'].unique()
# values = df['PaymentMethod'].value_counts()
#
# fig = go.Figure(data=[go.Pie(labels=labels, values=values, hole=.3)])
# fig.update_layout(title_text="<b>Payment Method Distribution</b>")
# fig.show()
```

Payment Method Distribution




Figure 4: Payment Method Distribution

```
# fig = px.histogram(df, x="Churn", color="PaymentMethod",
#                    title="<b>Customer Payment Method distribution w.r.t. Churn</b>")
# fig.update_layout(width=700, height=500, bargap=0.1)
# fig.show()
```




Figure 5: Customer Payment Method distribution w.r.t. Churn

Major customers who moved out were having Electronic Check as Payment Method.

Customers who opted for Credit-Card automatic transfer or Bank Automatic Transfer and Mailed Check as Payment Method were less likely to move out.

```
df["InternetService"].unique()
```

```
## array(['DSL', 'Fiber optic', 'No'], dtype=object)
```

```
df[df["gender"]=="Male"][["InternetService", "Churn"]].value_counts()
```

```
## InternetService  Churn
## DSL              No      992
## Fiber optic     No      910
## No              No      717
## Fiber optic     Yes      633
## DSL             Yes      240
## No             Yes       57
## dtype: int64
```

```
df[df["gender"]=="Female"][["InternetService", "Churn"]].value_counts()
```

```
## InternetService  Churn
## DSL              No      965
## Fiber optic     No      889
## No              No      690
## Fiber optic     Yes      664
## DSL             Yes      219
## No             Yes       56
## dtype: int64
```

A lot of customers choose the Fiber optic service and it's also evident that the customers who use Fiber optic have high churn rate, this might suggest a dissatisfaction with this type of internet service.

```

# fig = go.Figure()
# fig.add_trace(go.Bar(
#     x = [['Churn:No', 'Churn:No', 'Churn:Yes', 'Churn:Yes'],
#         ["Female", "Male", "Female", "Male"]],
#     y = [965, 992, 219, 240],
#     name = 'DSL',
# ))
# fig.add_trace(go.Bar(
#     x = [['Churn:No', 'Churn:No', 'Churn:Yes', 'Churn:Yes'],
#         ["Female", "Male", "Female", "Male"]],
#     y = [889, 910, 664, 633],
#     name = 'Fiber optic',
# ))
# fig.add_trace(go.Bar(
#     x = [['Churn:No', 'Churn:No', 'Churn:Yes', 'Churn:Yes'],
#         ["Female", "Male", "Female", "Male"]],
#     y = [690, 717, 56, 57],
#     name = 'No Internet',
# ))
# fig.update_layout(
#     title_text="Churn Distribution w.r.t. Internet Service and Gender")
# fig.show()

```




Figure 6: Churn Distribution w.r.t. Internet Service and Gender

```
# color_map = {"Yes": "#FF97FF", "No": "#AB63FA"}
# fig = px.histogram(df, x="Churn", color="Dependents",
#                    barmode="group", title="<b>Dependents distribution</b>",
#                    color_discrete_map=color_map)
# fig.update_layout(width=700, height=500, bargap=0.1)
# fig.show()
```

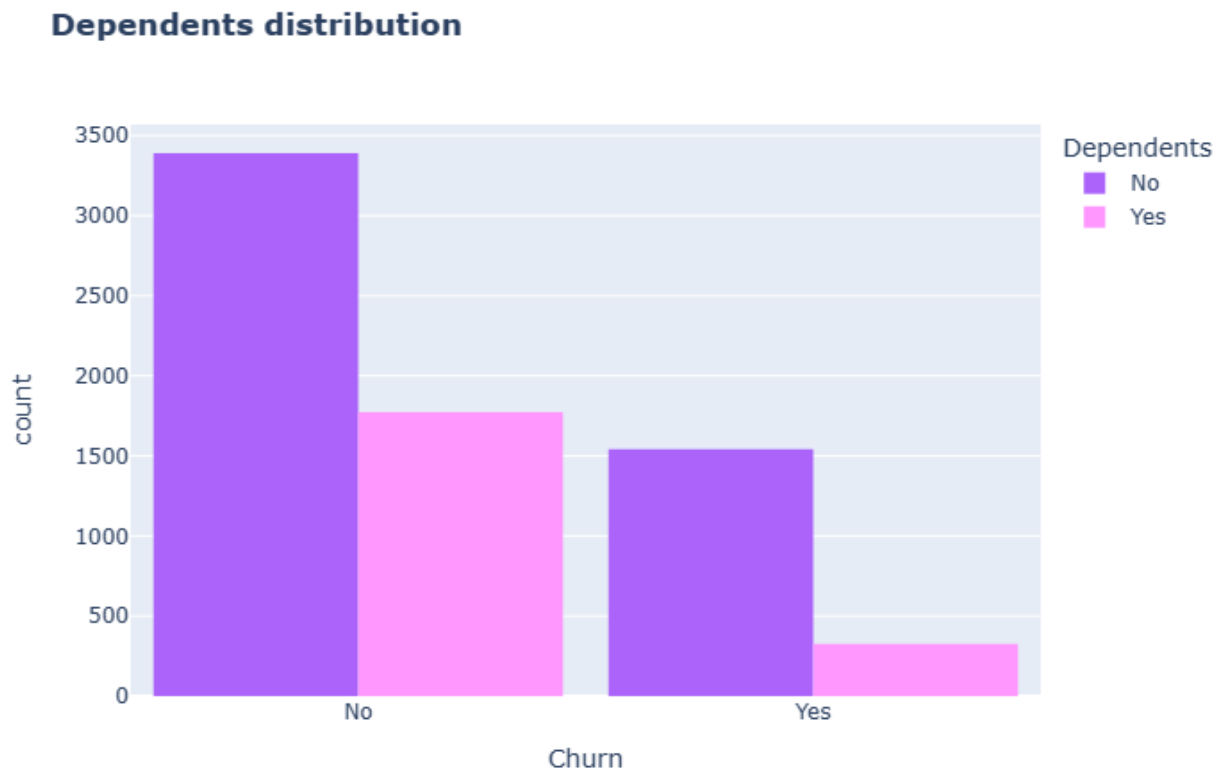


Figure 7: Dependents distribution

Customers without dependents are more likely to churn.

```
# color_map = {"Yes": '#FFA15A', "No": '#00CC96'}
# fig = px.histogram(df, x="Churn", color="Partner",
#                    barmode="group", title="Churn distribution w.r.t. Partners",
#                    color_discrete_map=color_map)
# fig.update_layout(width=700, height=500, bargap=0.1)
# fig.show()
```




Figure 8: Chrun distribution w.r.t. Partners

Customers that doesn't have partners are more likely to churn.


```
# color_map = {"Yes": '#00CC96', "No": '#B6E880'}
# fig = px.histogram(df, x="Churn", color="SeniorCitizen",
#                    title="<b>Chrun distribution w.r.t. Senior Citizen</b>",
#                    color_discrete_map=color_map)
# fig.update_layout(width=700, height=500, bargap=0.1)
# fig.show()
```

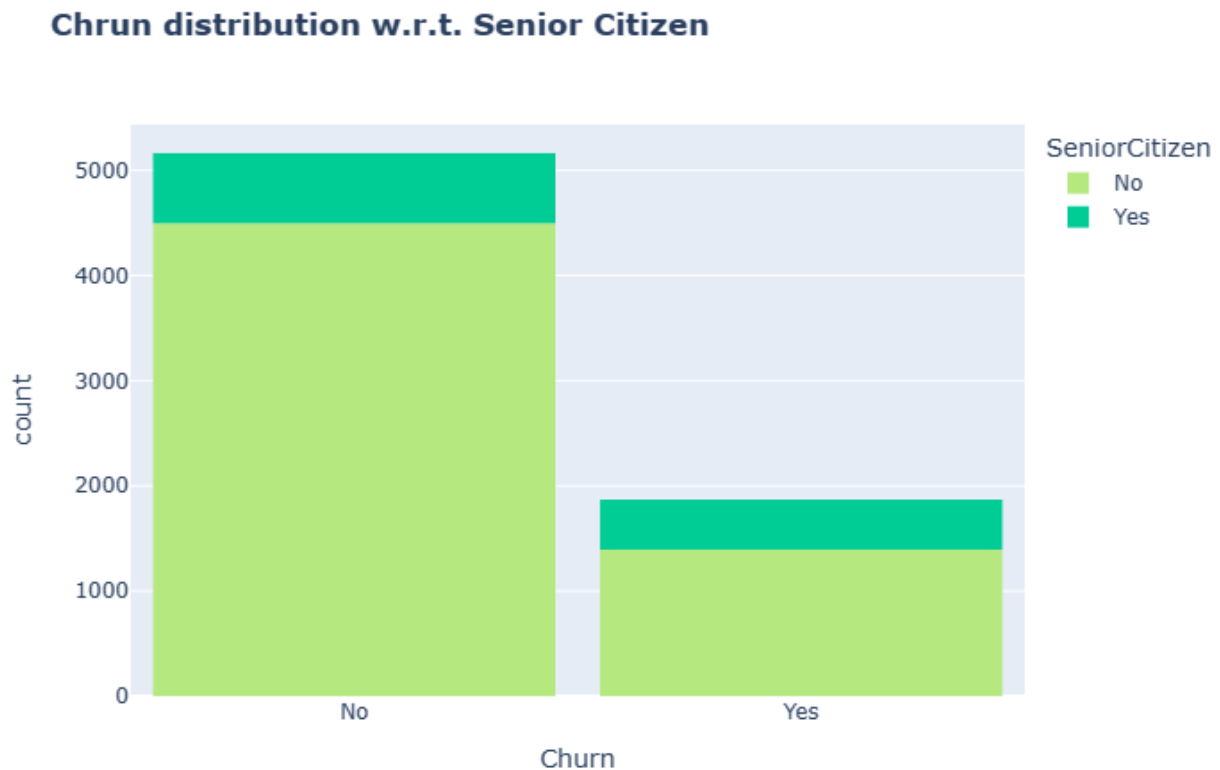


Figure 9: Chrun distribution w.r.t. Senior Citizen

It can be observed that the fraction of senior citizen is very less.
Most of the senior citizens do not churn.

```
# color_map = {"Yes": "#FF97FF", "No": "#AB63FA"}
# fig = px.histogram(df, x="Churn", color="OnlineSecurity",
#                    barmode="group", title="Churn w.r.t Online Security",
#                    color_discrete_map=color_map)
# fig.update_layout(width=700, height=500, bargap=0.1)
# fig.show()
```




Figure 10: Churn w.r.t Online Security

Most customers churn in the absence of online security.

```
# color_map = {"Yes": '#FFA15A', "No": '#00CC96'}
# fig = px.histogram(df, x="Churn", color="PaperlessBilling",
#                    title="<b>Chrun distribution w.r.t. Paperless Billing</b>",
#                    color_discrete_map=color_map)
# fig.update_layout(width=700, height=500, bargap=0.1)
# fig.show()
```

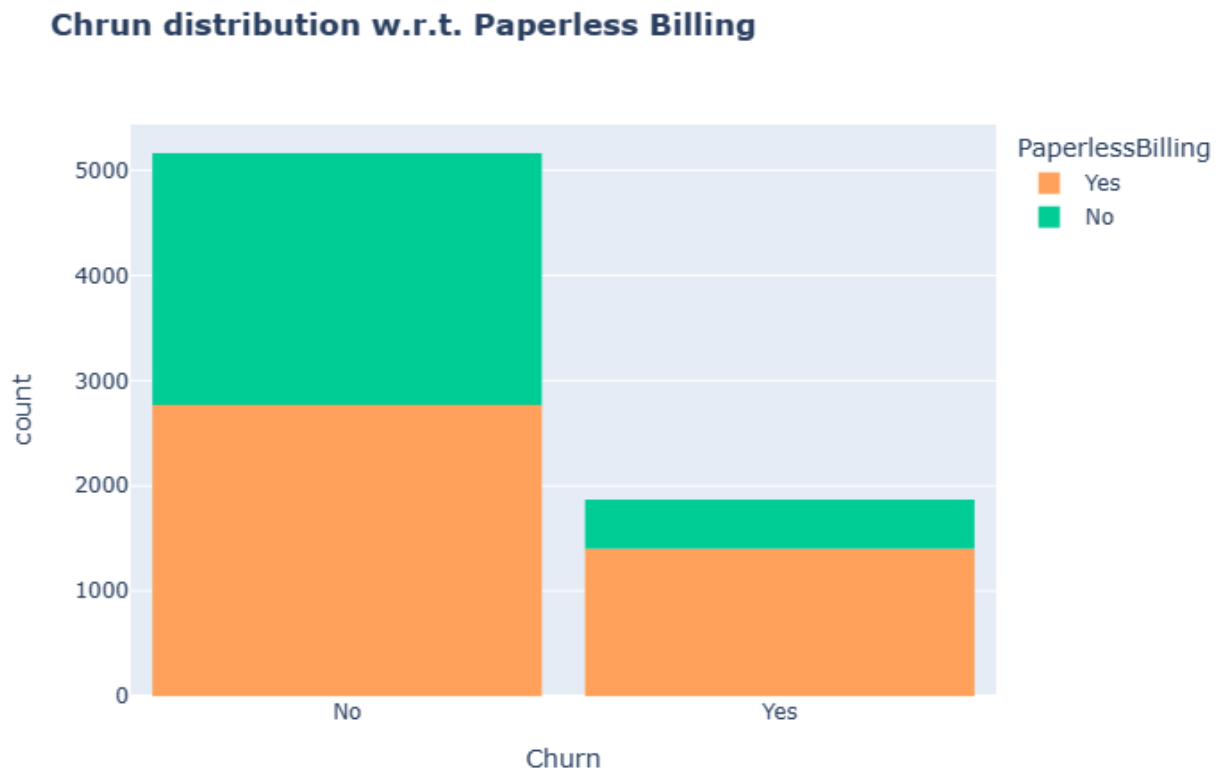


Figure 11: Chrun distribution w.r.t. Paperless Billing

Customers with Paperless Billing are most likely to churn.

```
# fig = px.histogram(df, x="Churn", color="TechSupport",
# barmode="group", title="Churn distribution w.r.t. TechSupport")
# fig.update_layout(width=700, height=500, bargap=0.1)
# fig.show()
```




Figure 12: Chrun distribution w.r.t. TechSupport

Customers without Tech Support are most likely to migrate to another service provider.

```
# color_map = {"Yes": '#00CC96', "No": '#B6E880'}
# fig = px.histogram(df, x="Churn", color="PhoneService",
#                    title="<b>Chrun distribution w.r.t. Phone Service</b>",
#                    color_discrete_map=color_map)
# fig.update_layout(width=700, height=500, bargap=0.1)
# fig.show()
```

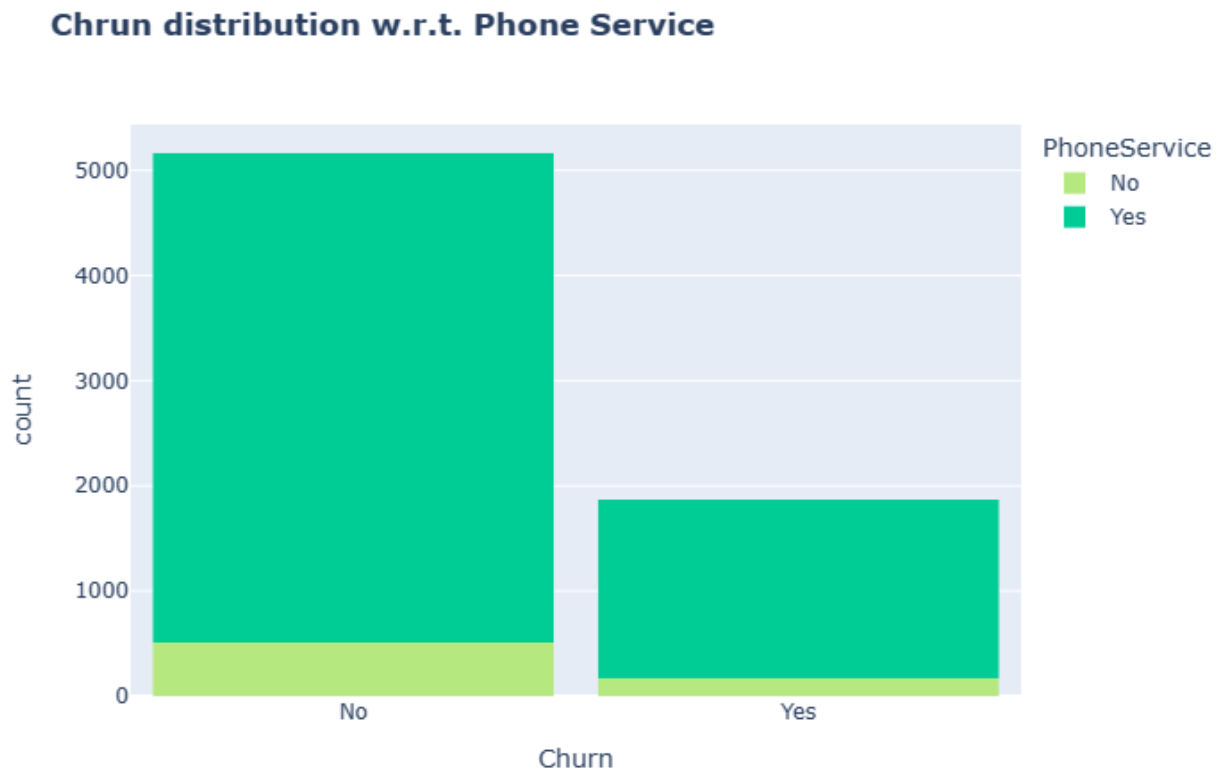


Figure 13: Chrun distribution w.r.t. Phone Service

Very small fraction of customers don't have a phone service.


Customers who have a phone service are more likely to churn.

```

no_churn = df[df["Churn"] == 'No']["MonthlyCharges"]
yes_churn = df[df["Churn"] == 'Yes']["MonthlyCharges"]

plt.figure(figsize=(10, 6))
sns.kdeplot(no_churn, color="red", shade=True, label="Not Churn")
sns.kdeplot(yes_churn, color="blue", shade=True, label="Churn")
plt.xlabel("Monthly Charges")
plt.ylabel("Density")
plt.title("Distribution of monthly charges by churn")
plt.legend(loc="upper right")
plt.show()

```




Customers with higher Monthly Charges are more likely to churn.

```

no_churn = df[df["Churn"] == 'No']["TotalCharges"]
yes_churn = df[df["Churn"] == 'Yes']["TotalCharges"]

plt.figure(figsize=(10, 6))
sns.kdeplot(no_churn, color="gold", shade=True, label="Not Churn")
sns.kdeplot(yes_churn, color="green", shade=True, label="Churn")
plt.xlabel("Total Charges")
plt.ylabel("Density")
plt.title("Distribution of total charges by churn")
plt.legend(loc="upper right")
plt.show()

```



Customers with lower Total Charges are more likely to churn.

```
# fig = px.box(df, x='Churn', y = 'tenure')
# # Update yaxis properties
# fig.update_yaxes(title_text='Tenure (Months)', row=1, col=1)
# # Update xaxis properties
# fig.update_xaxes(title_text='Churn', row=1, col=1)
# # Update size and title
# fig.update_layout(autosize=True, width=750, height=600,
#     title_font=dict(size=25, family='Courier'),
#     title='<b>Tenure vs Churn</b>')
# fig.show()
```

Tenure vs Churn




Figure 14: Tenure vs Churn

New customers are more likely to churn.


```

# Calculate the correlation matrix
corr = df.apply(lambda x: pd.factorize(x)[0]).corr()

# Filter the correlation matrix
# to include only correlations with the 'Churn' variable
churn_corr = pd.DataFrame(corr['Churn'])
churn_corr = churn_corr.reset_index()
churn_corr = churn_corr.drop(churn_corr.index[-1])

```

```

import matplotlib.pyplot as plt

# Generate a color map with unique colors
num_xticks = len(churn_corr)
color_map = plt.get_cmap('tab20')

# Plotting the correlation matrix
plt.figure(figsize=(10, 8))
bars = plt.bar(churn_corr['index'], churn_corr['Churn'])

# Assigning unique colors to xtick labels
for i, bar in enumerate(bars):
    bar.set_color(color_map(i % num_xticks))

# Set xtick label color to black
plt.xticks(rotation=45, color='black', fontsize=4)

```

```

plt.ylabel('Correlation with Churn', fontsize=8)
plt.title('Correlation of Variables with Churn')
plt.show()

```


7. Data Preprocessing

Splitting the data into train and test sets

```
def object_to_int(dataframe_series):  
    if dataframe_series.dtype=='object':  
        dataframe_series = LabelEncoder().fit_transform(dataframe_series)  
    return dataframe_series
```

```
df = df.apply(lambda x: object_to_int(x))  
df.head()
```

```
##      gender  SeniorCitizen  Partner  ...  MonthlyCharges  TotalCharges  Churn  
## 0         0             0         1  ...         29.85         29.85        0  
## 1         1             0         0  ...         56.95        1889.50        0  
## 2         1             0         0  ...         53.85         108.15        1  
## 3         1             0         0  ...         42.30        1840.75        0  
## 4         0             0         0  ...         70.70         151.65        1  
##  
## [5 rows x 20 columns]
```

```
plt.figure(figsize=(14,7))
df.corr()['Churn'].sort_values(ascending = False)
```


```
## Churn          1.000000
## MonthlyCharges  0.192858
## PaperlessBilling 0.191454
## SeniorCitizen  0.150541
## PaymentMethod  0.107852
## MultipleLines   0.038043
## PhoneService    0.011691
## gender          -0.008545
## StreamingTV     -0.036303
## StreamingMovies -0.038802
## InternetService -0.047097
## Partner         -0.149982
## Dependents      -0.163128
## DeviceProtection -0.177883
## OnlineBackup    -0.195290
## TotalCharges    -0.199484
## TechSupport     -0.282232
## OnlineSecurity  -0.289050
## tenure          -0.354049
## Contract        -0.396150
## Name: Churn, dtype: float64
```

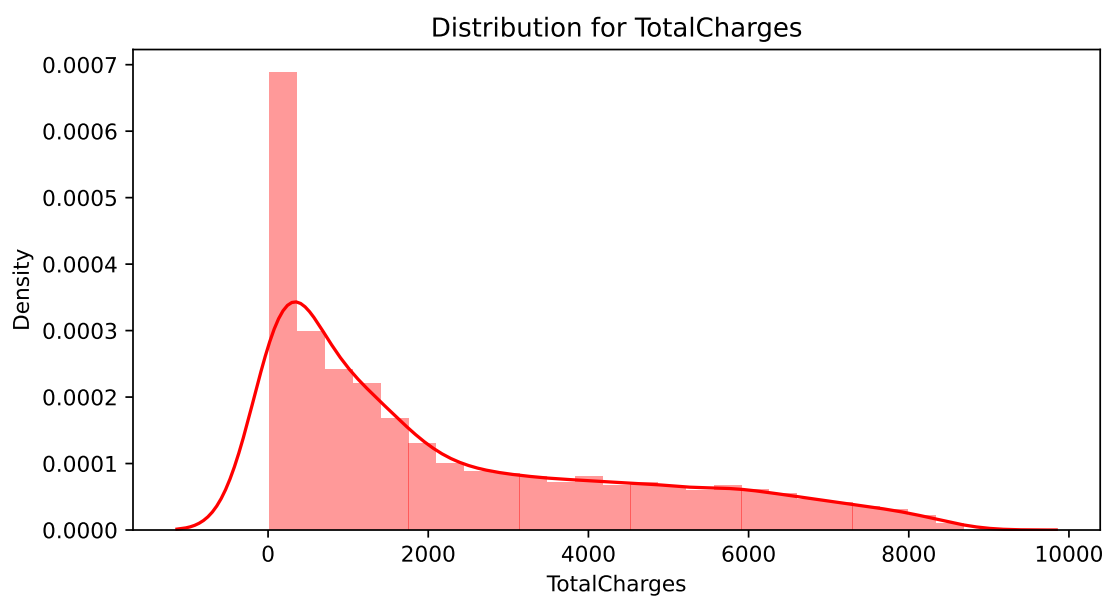

```
X = df.drop(columns = ['Churn'])
y = df['Churn'].values
```

```
X_train, X_test, y_train, y_test = train_test_split(X,y,
test_size = 0.30, random_state = 40, stratify=y)
```

```
def distplot(feature, frame, color='r'):
    plt.figure(figsize=(8,4))
    plt.title("Distribution for {}".format(feature))
    sns.distplot(frame[feature], color= color)
    plt.show()
```

```
num_cols = ["tenure", 'MonthlyCharges', 'TotalCharges']
for feat in num_cols: distplot(feat, df)
```








Since the numerical features are distributed over different value ranges, we will use standard scalar to scale them down to the same range.

Standardizing numeric attributes

```
df_std = pd.DataFrame(  
    StandardScaler().fit_transform(df[num_cols].astype('float64')),  
    columns=num_cols)  
  
for feat in numerical_cols: distplot(feat, df_std, color='c')
```





```
scaler = StandardScaler()

X_train[num_cols] = scaler.fit_transform(X_train[num_cols])
X_test[num_cols] = scaler.transform(X_test[num_cols])
```


8. Machine Learning Model Evaluations and Predictions

8.1 Logistic Regression

```
lr_model = LogisticRegression(random_state=3)
lr_model.fit(X_train,y_train)
```

```
## LogisticRegression(random_state=3)
```

```
prediction_lr = lr_model.predict(X_test)
accuracy_lr = accuracy_score(y_test, prediction_lr)
print("Logistic Regression accuracy is :", accuracy_lr)
```

```
## Logistic Regression accuracy is : 0.8090047393364929
```

8.2 AdaBoost

```
a_model = AdaBoostClassifier(random_state=3)
a_model.fit(X_train,y_train)
```

```
## AdaBoostClassifier(random_state=3)
```

```
prediction_a = a_model.predict(X_test)
accuracy_a = accuracy_score(y_test, prediction_a)
print("AdaBoost Classifier accuracy is :", accuracy_a)
```

```
## AdaBoost Classifier accuracy is : 0.8075829383886256
```

8.3 Gradient Boosting

```
gb_model = GradientBoostingClassifier(random_state=3)
gb_model.fit(X_train, y_train)
```

```
## GradientBoostingClassifier(random_state=3)
```

```
prediction_gb = gb_model.predict(X_test)
accuracy_gb = accuracy_score(y_test, prediction_gb)
print("Gradient Boosting Classifier accuracy is", accuracy_gb)
```

```
## Gradient Boosting Classifier accuracy is 0.8080568720379147
```

8.4 Voting Classifier

```
from sklearn.ensemble import VotingClassifier
lr = LogisticRegression(random_state=3)
abc = AdaBoostClassifier(random_state=3)
gbc = GradientBoostingClassifier(random_state=3)

eclf = VotingClassifier(estimators=[('lr', lr), ('abc', abc), ('gbc', gbc)],
                        voting='soft', weights=[1,1,1])

eclf.fit(X_train, y_train)
```

```
## VotingClassifier(estimators=[('lr', LogisticRegression(random_state=3)),
##                               ('abc', AdaBoostClassifier(random_state=3)),
##                               ('gbc',
##                                GradientBoostingClassifier(random_state=3))],
##                  voting='soft', weights=[1, 1, 1])
```

```
predictions = eclf.predict(X_test)
accuracy_vot = accuracy_score(y_test, predictions)
print("Voting Classifier Accuracy Score :", accuracy_vot)
```

```
## Voting Classifier Accuracy Score : 0.8170616113744076
```

8.5 Feature importance based on Voting Classifier Model

```
import eli5
from eli5.sklearn import PermutationImportance

perm = PermutationImportance(eclf, random_state=1).fit(X_test, y_test)
weights_df = eli5.formatters.as_dataframe.explain_weights_df(perm,
                                                             feature_names=X_test.columns.tolist())

print(weights_df)
```

##	feature	weight	std
## 0	tenure	0.075829	0.010283
## 1	MonthlyCharges	0.044645	0.006638
## 2	Contract	0.037820	0.003357
## 3	TotalCharges	0.012417	0.002671
## 4	TechSupport	0.011469	0.002654
## 5	PhoneService	0.007014	0.003357
## 6	OnlineSecurity	0.006730	0.002550
## 7	PaperlessBilling	0.005687	0.001038
## 8	SeniorCitizen	0.003223	0.001320
## 9	InternetService	0.002370	0.002952
## 10	PaymentMethod	0.002180	0.001546
## 11	MultipleLines	0.001991	0.001569
## 12	OnlineBackup	0.001706	0.001711
## 13	DeviceProtection	0.001422	0.001236
## 14	StreamingMovies	0.000853	0.000355
## 15	StreamingTV	0.000474	0.000599
## 16	gender	0.000095	0.000355
## 17	Dependents	-0.000569	0.001511
## 18	Partner	-0.000664	0.000643

According to this table, it's better to remove some features for the predictive modeling.

```

# Create new subsets of data with only the important features
columns_to_drop = ['StreamingMovies', 'StreamingTV', 'gender']

X_train_selected = X_train.drop(columns=columns_to_drop)
X_test_selected = X_test.drop(columns=columns_to_drop)

# Train the individual classifiers on the new data
lr_selected = LogisticRegression(random_state=3)
abc_selected = AdaBoostClassifier(random_state=3)
gbc_selected = GradientBoostingClassifier(random_state=3)

# Update the Voting Classifier with the new classifiers
ecf_selected = VotingClassifier(estimators=[('lr', lr_selected),
                                           ('abc', abc_selected), ('gbc', gbc_selected)],
                               voting='soft', weights=[1, 1, 1])

ecf_selected.fit(X_train_selected, y_train)

## VotingClassifier(estimators=[('lr', LogisticRegression(random_state=3)),
##                             ('abc', AdaBoostClassifier(random_state=3)),
##                             ('gbc',
##                               GradientBoostingClassifier(random_state=3))],
##                  voting='soft', weights=[1, 1, 1])

predictions_selected = ecf_selected.predict(X_test_selected)
accuracy_vot_selected = accuracy_score(y_test, predictions_selected)

print("Final Voting Classifier Accuracy Score:", accuracy_vot_selected)

## Final Voting Classifier Accuracy Score: 0.8175355450236966

```

9. Advice

After conducting feature selection by removing the attributes “Streaming Movies”, “Streaming TV” and “gender”, our refined machine learning model demonstrates an enhanced accuracy score of 81.75%. This noteworthy improvement surpasses the previous accuracy score of the Voting Classifier, which stood at 81.70%.

Based on the observed enhancement in our predictive model’s performance, it is advisable not to implement differential pricing based on gender. Additionally, it is recommended to exclude the features related to Streaming Movies and Streaming TV Services. By adhering to these recommendations, the company can optimize its customer retention strategies, thereby strengthening customer satisfaction and loyalty.